

Clayton Yochum

ML Systems · Performance · GPU Infrastructure

Salt Lake City, UT
claytonjy@gmail.com
linkedin.com/in/claytonjy
github.com/claytonjy

Machine learning systems engineer building and optimizing production ML services and GPU infrastructure. I establish technical foundations on small teams and scale them through company growth.

TECHNICAL SKILLS

ML & inference: PyTorch, NVIDIA Triton Inference Server, TensorRT-LLM, CTranslate2, NVIDIA Riva

Infrastructure: Kubernetes (GKE), GCP, AWS, Terraform, Helm, ArgoCD, KEDA, GitHub Actions

Services & data: Python, asyncio, FastAPI, Pydantic, Temporal, gRPC, WebSockets, SQL (PostgreSQL)

EXPERIENCE

Abridge

Senior Machine Learning Engineer

Machine Learning Engineer

Remote

Aug 2024 - Dec 2025

Apr 2023 - Jul 2024

Employee ~25 and ML engineer #2 at a Series B healthcare startup that grew to 400+ employees

- Designed and operated a multi-model speech-to-text (ASR) inference platform on GKE with NVIDIA Triton, processing **4 hours of audio per second** at daily peaks and up to 8 during backlog recovery.
- Migrated Whisper ASR from **2,000+ NVIDIA L4 GPUs** to **100-200 H100s** by replacing CTranslate2 with TensorRT-LLM and increasing batch sizes, serving the same production load at lower cost per audio-minute.
- Deployed real-time ASR on GKE and H100s using NVIDIA Parakeet and Riva, streaming in-app transcription with **sub-second end-to-end latency** over WebSockets and gRPC.
- Implemented a unified FastAPI interface across ML services, standardizing upstream integration, observability, and operational patterns.
- Introduced Temporal for complex asynchronous workflows and championed its adoption across nearly all engineering teams.
- Created monorepo development and release tooling for dozens of engineers, enabling cross-component changes in a single PR and reducing support requests for PRs and deployments.
- Mentored engineers in asynchronous Python and Helm deployment, enabling them to build and deploy new services independently.
- Led incident response for recurring networking issues and drove connection-reuse improvements across Python and JavaScript services.

Socure

Staff Data Engineer, Data Science Enablement & Operations

Remote

Jun 2022 - Jan 2023

Most senior engineer supporting data science teams with tooling and process

- Refined an internal batch API with service owners and shipped its Airflow integration.
- Developed reusable Airflow helpers around an autogenerated Python client for the customer-facing API.

Betterment

Lead Data Engineer

Remote

Jun 2021 - Dec 2021

Senior engineer on a centralized data engineering team supporting production databases and pipelines

- Standardized core Airflow reporting on Python, replacing R to simplify maintenance and ownership.
- Provided production on-call coverage for Airflow pipelines, databases, and Redshift replication, including overnight response and escalation to owning teams.

EXPERIENCE CONTINUED

Everactive

Senior Staff Engineer, Cloud Platform

Staff Engineer, Data Science

Ann Arbor, MI

Jan 2021 - May 2021

Mar 2019 - Jan 2021

First data science hire at a Series B industrial IoT startup

- Unified sensor readings and metadata in TimescaleDB/PostgreSQL, designing new schemas and migrating **billions of rows online** from OpenTSDB/EMR to simplify operations and accelerate queries.
- Flattened an overly nested sensor data model through further online migrations to improve TimescaleDB write performance.
- Served as technical lead for a company-wide API consolidation project, standardizing access to core sensor data across teams.
- Authored initial JSON Schema contracts between servers at customer sites and cloud services, securing buy-in from senior engineers on semantic versioning and compatibility requirements.
- Wrote job descriptions and conducted interviews for subsequent data science hires.

Methods Consultants

Data Science Engineer

Ypsilanti, MI

Mar 2016 - Feb 2019

First full-time hire and lead engineer at a boutique data science consulting firm

- Led development of an ML investment platform for the firm's largest client, training model ensembles to produce long/short beta-neutral portfolios directing **\$100k+ of daily trading volume**.
- Built a backtesting platform and automated model evaluation, backtesting, and daily prediction workflows on Pachyderm and GKE.
- Analyzed **billions of geolocation pings** with R and OpenStreetMap to support expert-witness testimony in a trucking fraud case.
- Developed and taught a week-long R and SQL workshop for analysts at a large federal institution.
- Replaced Dropbox code sharing with Git/GitLab and introduced CI/CD and AWS/GCP deployments.
- Hired and mentored junior data scientists.

Child and Adolescent Data Lab, University of Michigan

Data Analyst

Ann Arbor, MI

Aug 2015 - Dec 2017

Part-time industry consultant for a data-driven social work lab

- Built predictive models in R to classify journal abstracts, redesigned administrative data pipelines, and taught R, SQL, and Git across the School of Social Work.

Barracuda Networks

Business Intelligence Analyst

Performance Analyst/Software Engineer in Test

Ann Arbor, MI

Mar 2014 - Feb 2016

Aug 2012 - Mar 2014

QA engineer turned founding business intelligence analyst at a large software company

- Built a MySQL warehouse with PHP ingestion jobs to record nightly backup-server utilization across regions.
- Developed R Shiny dashboards to forecast regional storage growth and inform hardware purchasing decisions.
- Wrote parallel R jobs using MPI to calculate deduplication efficiency for a global backup service, overcoming memory constraints in long-running analyses.

EDUCATION

Michigan Technological University

Bachelor of Science in Mathematics, minor in Economics

Houghton, MI

2012