

Clayton Yochum

ML Systems · Inference · GPU Infrastructure

Salt Lake City, UT
claytonjy@gmail.com
linkedin.com/in/claytonjy
github.com/claytonjy

ML Systems Engineer with 15 years experience across Machine Learning, Data Engineering, and Data Science. I build the infrastructure between models and products: inference APIs, GPU orchestration, and production reliability. I like wearing many hats on small teams, building early foundations and scaling systems as the company grows.

TECHNICAL SKILLS

ML & inference: Python, PyTorch, speech recognition (ASR), TensorRT-LLM, CTranslate2, NVIDIA Triton; NVIDIA L4 and H100 GPUs.

Cloud & production systems: GCP, GKE/Kubernetes, AWS/ECS, Terraform, Helm, ArgoCD, KEDA, Temporal; incident investigation and deployment architecture.

Data & developer platforms: FastAPI, Pydantic, JSON Schema, Kafka, Airflow, PostgreSQL/TimescaleDB, SQL, R, GitHub Actions, GitLab, Pachyderm.

EXPERIENCE

Abridge

Senior Machine Learning Engineer
Machine Learning Engineer

Remote
Aug 2024 - Dec 2025
Apr 2023 - Jul 2024

Joined at Series B as employee 25 and ML engineer 2 scaled to 400+ employees and \$5B valuation

- Built and operated the in-house speech-recognition platform on GKE, transcribing **4-8 hours of audio per second** in production.
- Migrated ASR inference from **2,000+ NVIDIA L4 GPUs** running CTranslate2 to **100-200 H100s** running TensorRT-LLM for the same production traffic load.
- Designed a unified FastAPI interface across ML services, giving upstream teams one API for ML capabilities and centralizing observability and operational patterns.
- Introduced Temporal for complex asynchronous workflows, extending adoption beyond ML to other engineering teams.
- Partnered with Platform to establish engineering-wide infrastructure and deployment practices on GKE using Terraform, Helm, ArgoCD, and KEDA.
- Built GitHub Actions and release tooling for the ML monorepo, enabling cross-component changes in one PR and automating semantic versioning for dozens of contributors.
- Shaped architecture across ML and Platform; mentored engineers in development and deployment practices, influenced cross-team technical decisions, and conducted technical interviews with Platform.
- Led major production incident investigations across application, Kubernetes, networking, and infrastructure layers, including incident response and post-mortems for severe outages.

Socure

Staff Data Engineer, Data Science Enablement & Operations

Remote
Jun 2022 - Jan 2023

Supported Data Science, bridging Machine Learning and Data Platform teams.

- Integrated an internal service into the primary pipeline to reduce customer onboarding time.
- Wrote Python libraries to simplify Airflow pipeline development and migrated existing pipelines to new patterns.

Betterment

Lead Data Engineer

Remote
Jun 2021 - Dec 2021

- Rewrote core reporting pipelines in Airflow from R to Python.
- Introduced Pydantic for AWS Lambda tasks, and simplified the data engineering platform's developer experience to increase contributions from other teams.

EXPERIENCE CONTINUED

Everactive

Senior Staff Engineer, Cloud Platform

Staff Engineer, Data Science

Ann Arbor, MI

Jan 2021 - May 2021

Mar 2019 - Jan 2021

First Data Science hire at a Series B industrial IoT startup; primary data expert on a small Cloud Platform team.

- Led migration of **billions of rows of core sensor data** from OpenTSDB to TimescaleDB/PostgreSQL, simplifying the stack and improving customer-facing responsiveness.
- Served as technical lead for a company-wide API consolidation project, standardizing access to core sensor data across teams.
- Negotiated strict JSON Schema contracts with hardware teams to define interfaces for sensor data.
- Led the shift toward event-driven platform architecture on Kafka.
- Created Python libraries and AWS ECS services for automated failure detection from sensor data.
- Wrote job descriptions and conducted interviews for subsequent Data Science hires.

Methods Consultants

Data Science Engineer

Ypsilanti, MI

Mar 2016 - Feb 2019

First full-time hire and lead engineer at a boutique data science and statistics consulting firm.

- Developed a semi-automated ML trading platform in R and Python for a mid-size hedge fund, running on GKE and Pachyderm and supporting **over \$100k in daily trading volume**.
- Analyzed **billions of geolocation pings** with R and OpenStreetMap to support expert-witness testimony in a trucking fraud case.
- Developed and taught a **week-long R and SQL workshop** for analysts at a large federal institution.
- Introduced version control with Git and GitLab and cloud computing with AWS and GCP, establishing foundations for the firm's engineering work.
- Hired and mentored junior data scientists; contributed technical writing to the company blog, including A Tour of Timezones (& Troubles) in R.

University of Michigan School of Social Work

Data Analyst (part-time)

Ann Arbor, MI

Aug 2015 - Dec 2017

- Built predictive models to classify journal abstracts, redesigned pipelines for administrative data, and created reusable lab tooling.
- Taught source control and reproducible research to analysts and the broader School of Social Work.

Barracuda Networks

Business Intelligence Analyst

Performance Analyst/Software Engineer in Test

Ann Arbor, MI

Mar 2014 - Feb 2016

Aug 2012 - Mar 2014

- Founded and technically led the Analytics team
- Created data warehouses to provide our first historical view of critical production data.
- Managed server infrastructure and developed R tooling for cloud capacity planning.

EDUCATION

Michigan Technological University

Bachelor of Science in Mathematics, minor in Economics

2012